

```

from google.colab import files
import zipfile, os

# Upload your MLPColab.zip manually from local machine
uploaded = files.upload() # choose MLPColab.zip

# Extract
zip_path = "MLPColab.zip"
extract_dir = "/content/MLPColab"
os.makedirs(extract_dir, exist_ok=True)

with zipfile.ZipFile(zip_path, "r") as zf:
    zf.extractall(extract_dir)

os.chdir(extract_dir)
print("✅ Files extracted to:", os.getcwd())
!ls -R | head -40

```

No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
 Saving MLPColab.zip to MLPColab.zip
 Files extracted to: /content/MLPColab
 .:
 MLPColab

 ./MLPColab:
 data_prep.py
 dataset.py
 README.md
 requirements.txt
 streamlit_app.py
 train_mlp.py
 utils.py
 visualize_embeddings.py

```
!pip install torch torchvision tqdm streamlit pyngrok
```

```

Requirement already satisfied: torch in /usr/local/lib/python3.12/dist-p
Requirement already satisfied: torchvision in /usr/local/lib/python3.12/
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-pa
Collecting streamlit
  Downloading streamlit-1.51.0-py3-none-any.whl.metadata (9.5 kB)
Collecting pyngrok
  Downloading pyngrok-7.4.1-py3-none-any.whl.metadata (8.1 kB)
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dis
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/]
Requirement already satisfied: setuptools in /usr/local/lib/python3.12/c
Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3.1
Requirement already satisfied: networkx in /usr/local/lib/python3.12/dis

```

```
Requirement already satisfied: jinja2 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: fsspec in /usr/local/lib/python3.12/dist-
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.6.77 in /usr/>
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.6.77 in /usr>
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.6.80 in /usr/>
Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in /usr/loc
Requirement already satisfied: nvidia-cublas-cu12==12.6.4.1 in /usr/loc
Requirement already satisfied: nvidia-cufft-cu12==11.3.0.4 in /usr/local
Requirement already satisfied: nvidia-curand-cu12==10.3.7.77 in /usr/loc
Requirement already satisfied: nvidia-cusolver-cu12==11.7.1.2 in /usr/loc
Requirement already satisfied: nvidia-cusparse-cu12==12.5.4.2 in /usr/loc
Requirement already satisfied: nvidia-cusparseelt-cu12==0.7.1 in /usr/loc
Requirement already satisfied: nvidia-nccl-cu12==2.27.3 in /usr/local/li
Requirement already satisfied: nvidia-nvtx-cu12==12.6.77 in /usr/local/]
Requirement already satisfied: nvidia-nvjitlink-cu12==12.6.85 in /usr/loc
Requirement already satisfied: nvidia-cufile-cu12==1.11.1.6 in /usr/loc
Requirement already satisfied: triton==3.4.0 in /usr/local/lib/python3.1
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-p
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /usr/local/lib/p
Requirement already satisfied: altair!=5.4.0,!<5.4.1,<6,>=4.0 in /usr/lc
Requirement already satisfied: blinker<2,>=1.5.0 in /usr/local/lib/pythc
Requirement already satisfied: cachetools<7,>=4.0 in /usr/local/lib/pyth
Requirement already satisfied: click<9,>=7.0 in /usr/local/lib/python3.1
Requirement already satisfied: packaging<26,>=20 in /usr/local/lib/pythc
Requirement already satisfied: pandas<3,>=1.4.0 in /usr/local/lib/pythor
Requirement already satisfied: protobuf<7,>=3.20 in /usr/local/lib/pythc
Requirement already satisfied: pyarrow<22,>=7.0 in /usr/local/lib/pythor
Requirement already satisfied: requests<3,>=2.27 in /usr/local/lib/pythc
Requirement already satisfied: tenacity<10,>=8.1.0 in /usr/local/lib/pyt
Requirement already satisfied: toml<2,>=0.10.1 in /usr/local/lib/python3
Requirement already satisfied: watchdog<7,>=2.1.5 in /usr/local/lib/pyth
Requirement already satisfied: gitpython!=3.1.19,<4,>=3.0.7 in /usr/loc
Collecting pydeck<1,>=0.8.0b4 (from streamlit)
  Downloading pydeck-0.9.1-py2.py3-none-any.whl.metadata (4.1 kB)
Requirement already satisfied: tornado!=6.5.0,<7,>=6.0.3 in /usr/local/]
Requirement already satisfied: PyYAML>=5.1 in /usr/local/lib/python3.12/
Requirement already satisfied: jsonschema>=3.0 in /usr/local/lib/python3
Requirement already satisfied: narwhals>=1.14.2 in /usr/local/lib/pythor
Requirement already satisfied: gitdb<5,>=4.0.1 in /usr/local/lib/python3
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3
Requirement already satisfied: charset_normalizer<4,>=2 in /usr/local/li
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.12
```

```
import os, json, requests, numpy as np
from tqdm import tqdm

base_data = "data"
os.makedirs(base_data, exist_ok=True)

datasets = {
    "warpeace": "https://raw.githubusercontent.com/karpathy/char-rnn/master/datasets/warpeace.txt",
    "linux": "https://raw.githubusercontent.com/karpathy/char-rnn/master/datasets/linux.txt"
}
```

```

for name, url in datasets.items():
    ddir = os.path.join(base_data, name)
    os.makedirs(ddir, exist_ok=True)

    txt_path = os.path.join(ddir, "input.txt")
    if not os.path.exists(txt_path):
        print(f"⬇️ Downloading {name} dataset...")
        txt = requests.get(url).text
        open(txt_path, "w", encoding="utf-8").write(txt)

    # Build vocab and encoded files (robust fallback)
    text = open(txt_path, "r", encoding="utf-8").read()
    vocab = sorted(list(set(text)))
    stoi = {ch:i for i,ch in enumerate(vocab)}
    itos = {i:ch for ch,i in stoi.items()}
    json.dump({"stoi": stoi, "itos": itos}, open(os.path.join(ddir, "vocab.json"), "w"))

    data = np.array([stoi[ch] for ch in text], dtype=np.uint16)
    split = int(0.9 * len(data))
    np.save(os.path.join(ddir, "train.npy"), data[:split])
    np.save(os.path.join(ddir, "val.npy"), data[split:])
    print("✅ Data ready.")

```

Downloading warpeace dataset...
 Downloading linux dataset...
 Data ready.

```

import os, json, numpy as np
from tqdm import tqdm

base_data = "data"

for name in ["warpeace", "linux"]:
    ddir = os.path.join(base_data, name)
    vocab_json = os.path.join(ddir, "vocab.json")
    raw_txt = os.path.join(ddir, f"{name}.txt")
    tokens_file = os.path.join(ddir, "tokens.ids")

    if not os.path.exists(vocab_json):
        raise FileNotFoundError(f"Missing {vocab_json}")
    if not os.path.exists(raw_txt):
        raise FileNotFoundError(f"Missing {raw_txt}")

    info = json.load(open(vocab_json, "r", encoding="utf-8"))
    stoi = info["stoi"]

    print(f"✍️ Generating tokens.ids for {name} ...")
    with open(raw_txt, "r", encoding="utf-8") as f:
        text = f.read()

    # Simple tokenization (char-level)
    tokens = [stoi[ch] for ch in tqdm(text, desc=f"Tokenizing {name}") if ch in stoi]

    np.savetxt(tokens_file, tokens, fmt="%d")

```

```
print(f"✅ Saved {len(tokens)} tokens → {tokens_file}")

-----
---  
FileNotFoundError                         Traceback (most recent call
last)
/tmp/ipython-input-2151976901.py in <cell line: 0>()
    13         raise FileNotFoundError(f"Missing {vocab_json}")
    14     if not os.path.exists(raw_txt):
---> 15         raise FileNotFoundError(f"Missing {raw_txt}")
    16
    17     info = json.load(open(vocab_json, "r", encoding="utf-8"))

FileNotFoundError: Missing data/warpeace/warpeace.txt
```

```
!python train_mlp.py --dataset warpeace --epochs 30 --batch_size 64 --lr 1e-3
!python train_mlp.py --dataset linux --epochs 30 --batch_size 64 --lr 1e-3 --s
```

```
Device: cpu
Traceback (most recent call last):
  File "/content/MLPColab/MLPColab/train_mlp.py", line 131, in <module>
    main(args)
  File "/content/MLPColab/MLPColab/train_mlp.py", line 52, in main
    train_ds = NextWordDataset(args.dataset, context_len=args.context_len
                                ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "/content/MLPColab/MLPColab/dataset.py", line 20, in __init__
    with open(tokens_path, "r", encoding="utf-8") as f:
                                ^^^^^^^^^^^^^^
FileNotFoundException: [Errno 2] No such file or directory: 'data/warpeace/to
Device: cpu
Traceback (most recent call last):
  File "/content/MLPColab/MLPColab/train_mlp.py", line 131, in <module>
    main(args)
  File "/content/MLPColab/MLPColab/train_mlp.py", line 52, in main
    train_ds = NextWordDataset(args.dataset, context_len=args.context_len
                                ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "/content/MLPColab/MLPColab/dataset.py", line 20, in __init__
    with open(tokens_path, "r", encoding="utf-8") as f:
                                ^^^^^^^^^^^^^^
FileNotFoundException: [Errno 2] No such file or directory: 'data/linux/token
```

```
from pyngrok import ngrok
import subprocess, time

# Kill old sessions
!fuser -k 8501/tcp || echo "No old Streamlit session"

# Start Streamlit (must exist in same dir)
```

```

process = subprocess.Popen(["streamlit", "run", "app.py", "--server.port", "8501"])
time.sleep(6)

# Connect to ngrok tunnel
public_url = ngrok.connect(8501)
print("✅ Streamlit running at:", public_url)

```

```

# LINUX training (30 epochs) - smaller batch to avoid OOM
python train_mlp.py \
--dataset linux \
--context_len 5 \
--embed_dim 64 \
--hidden_size 512 \
--n_hidden 1 \
--activation relu \
--dropout 0.2 \
--batch_size 64 \
--epochs 30 \
--lr 1e-3 \
--save_dir checkpoints/linux \
--save_every 5 \
--val_ratio 0.1

```

```

echo "War & Peace checkpoints:"
ls -la checkpoints/warpeace || true
echo "Linux checkpoints:"
ls -la checkpoints/linux || true

```

Warpeace

```

# determine model file (best.pt preferred)
if [ -f checkpoints/warpeace/best.pt ]; then MODEL_W=checkpoints/warpeace/best.pt
echo "Using War model: $MODEL_W"
python visualize_embeddings.py --dataset warpeace --model_path "$MODEL_W" --top_k 5

```

Linux

```

if [ -f checkpoints/linux/best.pt ]; then MODEL_L=checkpoints/linux/best.pt; else MODEL_L=checkpoints/warpeace/best.pt; fi
echo "Using Linux model: $MODEL_L"
python visualize_embeddings.py --dataset linux --model_path "$MODEL_L" --top_k 5

from IPython.display import Image, display
print("War & Peace embedding viz:")
display(Image('checkpoints/warpeace/warpeace_tsne.png'))
print("Linux embedding viz:")
display(Image('checkpoints/linux/linux_tsne.png'))

```

Optional (below)

```
%bash
python - <<'PY'
import torch, json, sys
from pathlib import Path

def load_ckpt(p):
    ck = torch.load(p, map_location='cpu')
    return ck

def simple_generate(ckpt_path, seed_text="the", k=30, context_len=5, temperature=1.0):
    ck = load_ckpt(ckpt_path)
    vocab = ck['vocab']
    args = ck.get('args', {})
    # Build model skeleton exactly like train_mlp MLPNextWord
    import torch.nn as nn
    class MLPNextWord(nn.Module):
        def __init__(self, vocab_size, context_len=5, embed_dim=64, hidden_size=64):
            super().__init__()
            self.embedding = nn.Embedding(vocab_size, embed_dim, padding_idx=0)
            layers = []
            in_dim = context_len * embed_dim
            for _ in range(n_hidden):
                layers.append(nn.Linear(in_dim, hidden_size))
                layers.append(nn.ReLU() if activation=="relu" else nn.Tanh())
                in_dim = hidden_size
            layers.append(nn.Linear(in_dim, vocab_size))
            self.net = nn.Sequential(*layers)
        def forward(self, x):
            e = self.embedding(x)
            flat = e.view(e.size(0), -1)
            return self.net(flat)
    margs = args
    model = MLPNextWord(vocab_size=len(vocab), context_len=margs.get('context_len'))
    model.load_state_dict(ck['model_state'], strict=False)
    model.eval()
    word2idx = {w:i for i,w in enumerate(vocab)}
    idx2word = {i:w for i,w in enumerate(vocab)}
    tokens = seed_text.lower().split()
    ids = [word2idx.get(t, word2idx.get("<UNK>",0)) for t in tokens]
    out = tokens[:]
    import numpy as np
    for _ in range(k):
        context = ids[-context_len:] if len(ids)>=context_len else [word2idx.get(t, word2idx.get("<UNK>",0)) for t in tokens[-context_len:]]
        x = torch.tensor([context], dtype=torch.long)
        logits = model(x)[0].detach().numpy()
        # temperature sampling
        if temperature<=0:
            nid = int(logits.argmax())
        else:
            probs = np.exp(logits/temperature)
            probs = probs / probs.sum()
            nid = int(np.random.choice(len(probs), p=probs))
        out.append(idx2word.get(nid, "<UNK>"))
        ids.append(nid)
    return " ".join(out)
```

```
# adjust these paths as appropriate
war_ckpt = Path("checkpoints/warpeace/best.pt")
if not war_ckpt.exists():
    war_ckpt = Path("checkpoints/warpeace/final.pt")
linux_ckpt = Path("checkpoints/linux/best.pt")
if not linux_ckpt.exists():
    linux_ckpt = Path("checkpoints/linux/final.pt")

print("War sample:")
print(simple_generate(str(war_ckpt), seed_text="the prince", k=30, context_len=100))
print("\nLinux sample:")
print(simple_generate(str(linux_ckpt), seed_text="static inline", k=30, context_len=100))
```

.app py created

```

        layers.append(nn.Linear(in_dim, hidden_size))
        layers.append(nn.ReLU() if activation=="relu" else nn.Tanh())
        in_dim = hidden_size
    layers.append(nn.Linear(in_dim, vocab_size))
    self.net = nn.Sequential(*layers)
def forward(self, x):
    e = self.embedding(x)
    flat = e.view(e.size(0), -1)
    return self.net(flat)
m = MLPNextWord(vocab_size=len(vocab), context_len=args.get('context_len',
m.load_state_dict(ckpt['model_state'], strict=False)
m.eval()
return m, vocab

if st.button("Load model & generate"):
    if not os.path.exists(model_file):
        st.error("Model file not found. Please correct model path.")
    else:
        ckpt = load_ckpt(model_file)
        model, vocab = build_model_from_ckpt(ckpt)
        w2i = {w:i for i,w in enumerate(vocab)}
        i2w = {i:w for i,w in enumerate(vocab)}
        tokens = seed_text.strip().lower().split()
        ids = [w2i.get(t, w2i.get("<UNK>",0)) for t in tokens]
        out = tokens[:]
        import numpy as np
        for _ in range(k):
            context = ids[-context_len:] if len(ids)>=context_len else [w2i.get(
            x = torch.tensor([context], dtype=torch.long)
            logits = model(x)[0].detach().numpy()
            if temperature<=0:
                nid = int(logits.argmax())
            else:
                probs = np.exp(logits/temperature)
                probs = probs / probs.sum()
                nid = int(np.random.choice(len(probs), p=probs))
            out.append(i2w.get(nid, "<UNK>"))
            ids.append(nid)
        st.subheader("Generated text")
        st.write(" ".join(out))

```

```

!./ngrok authtoken YOUR_NGROK_AUTHTOKEN

!kill $(lsof -t -i:8501) >/dev/null 2>&1 || echo "no process"

# run streamlit in background
import subprocess, time, os
proc = subprocess.Popen(["streamlit", "run", "app.py", "--server.port", "8501"])
time.sleep(6)

# create tunnel and print URL
from pyngrok import ngrok
public_url = ngrok.connect(8501)

```

```
print("Streamlit public URL:", public_url)
```